



Principles of Concurrent Systems

Processes and Threads

Designing Multi-tasking systems

- In comparison to a **single tasking** system the design of a **concurrent systems** represents a problem at least an **order of magnitude more complex** both in conceptual design (in terms of process design, interaction and communication), and in the implementation via a multi-tasking kernel.

Multi-Tasking Tools

- To successfully implement a concurrent real time system, the designer must have available a range of tools to support such concepts as:-
 - Process/thread Creation, Priority and Scheduling.
 - Process Communication
 - Process Synchronisation
 - Process Stimulation
 - Process Interaction with External asynchronous IO and Event

- These tools will probably include:-
 - A **multi-tasking kernel** (MTOS) to implement these concepts and provide a suitable environment under which these processes can be run e.g. Win32 or Linux.
 - A suitable **programming language** to support the ideas and requirements of concurrent real-time systems, providing the necessary interfaces to the real-time kernel to allow their realisation within user programs.
 - A **design tool** (C.A.S.E. Tool) to allow the designer to capture and express easily, in a graphical manner, the design and architecture of a complex, large scale multi-tasking systems. Ideally such tools should provide **consistency checking** and **database** support to facilitate the creation of re-usable components across projects.
- We shall look in turn at each of these requirements, after we have analysed in detail the problems associated with concurrency.

Problems of Concurrent Systems

- The fundamental difference between a single tasking system and a concurrent one is obviously the inclusion of several parallel **co-operating processes or threads**.
- The inclusion of multiple processes is of prime importance in the design. All other problems concerning concurrent systems stem directly from the inclusion of these processes, e.g. process **communication**, process **synchronisation**, and **stimulation**.
- Such problems only **exist** because of these multiple processes.

Design Methodologies for Concurrent Systems

- Many of the design approaches that have traditionally been applied to the design of single tasking systems can equally be applied to concurrent systems. Design ideas such as:-
 - Object Oriented Analysis and Design
 - Data Flow Analysis and Structured Design
 - State Behaviour Analysis e.g. state transition diagrams and charts.
 - Modular Programming
 - Top down and Hierarchical design
- However, because we now wish to design several processes instead of just one, there is an additional design concept called.

Hierarchical **Process** Decomposition

- Here, a large system is **broken down** into a collection of **processes** and **threads** which **co-operate** and **communicate** with each other in a variety of ways.
- Once the system has been expressed in terms of its processes/threads, the traditional design approaches mentioned above can be applied to the design of each process/thread in turn.

Introduction to Parallel Programming Concepts

- If a **concurrent** system is to be designed, there must exist a way of **creating** and **controlling** the parallel **activities** or **processes** within the system.
- Ideally we would like to be able to express naturally which parts of our programs we want to run in **parallel** and which parts we want to run **sequentially**.
- Ultimately such concurrency will actually be implemented on a single CPU using **time slicing** techniques controlled by a **real-time clock** interrupting the CPU at regular intervals. Here the kernel takes responsibility for '**swapping out**' one process and '**swapping in**' another.
- Any **programming language** that claims to support **concurrency** must do so by making hidden calls to the kernel which ideally are **transparent** to the programmer.
- In other words, the **compiler** translates the **concurrent programming expressions** in our code into the appropriate **calls to the kernel** to create and control multiple processes and threads.
- A good parallel programming language should make this easy for the programmer without them having any regard for the underlying CPU or host operating system

Parallel Programming Languages

- Some programming languages like **Java**, **Occam** and **Ada** were designed from the outset to support concurrent programming concepts with additional **key words** to the language to express these ideas.
- Other languages like '**C/C++**', **Fortran** etc. did not and special versions of those languages such as '**Concurrent C**' or **Concurrent Fortran** have existed for a number of years with extensions to the 'base' language to make parallel programming more natural

Concurrent Programming Language Extensions - **PAR** and **SEQ**

- Having a language with built in parallel programming extensions greatly simplifies the design of concurrent systems
- Most concurrent programming languages use key words such as '**PAR**' (parallel) and '**SEQ**' (sequence) to express parallel and sequential sections of a programs source code.
- For example in **Occam** we could write something like this within one source file.

```
PAR
  SEQ
    ...
  SEQ
    ...
  SEQ
    ...
PAR END
```

```
/* Process 1 */
/* Program code for process 1 */

/* Process 2 */
/* Program code for process 2 */

/* Process 3 */
/* Program code for process 3 */
```

- This of course is just a simple example, in reality, in a more complex system, the activity's done in parallel could be very complex.

Implementing Concurrency with Multiple Processes

- For those languages that **do not directly support concurrency**, such as **C/C++**, the chosen language must provide a **library** that your programs can use to make the direct calls to the kernel to achieve things like **creating threads** and **processes**.
- In reality this is what operations such as **PAR** and **SEQ** do, but they do it in a more user friendly manner that is **transparent** to the programmer and **independent** of the host operating system, i.e. the code should be portable, you just need a compiler to target the host operating system.
- Obviously placing such **explicit kernel code** into your programs is a less attractive proposition, since it effectively ties down the source code you have written to a specific kernel thus affecting the portability of the code to other environments, i.e. you cannot just recompile it for another operating system as the same kernel calls may not exist.
- Typically, this is not such a big problem as it might at first appear, since many kernels operate in much the same way providing many of the same features.
- Additionally, real-time applications are generally **heavily tied down** to the particular **hardware/platform** that they are running on and thus it is unlikely that it will ever be 'ported' to a radically different environment.
- However to make life easier, you can **wrap up the detailed and complex system calls** into a library which provide a greater level of programming **abstraction** to effectively hide away the explicit operating system code so that the **application code** becomes more portable, that is, you only have to rewrite the library for a different operating system and the whole application should port over.
- This effectively is what the **rt.cpp** source file is. It is a **high level wrapper** around the **Win32 Kernel**. Applications call code within **rt.cpp** to achieve a high level operation and leave the library to make the complex and detailed low level calls the operating system.

Process Creation: The **CProcess()** class

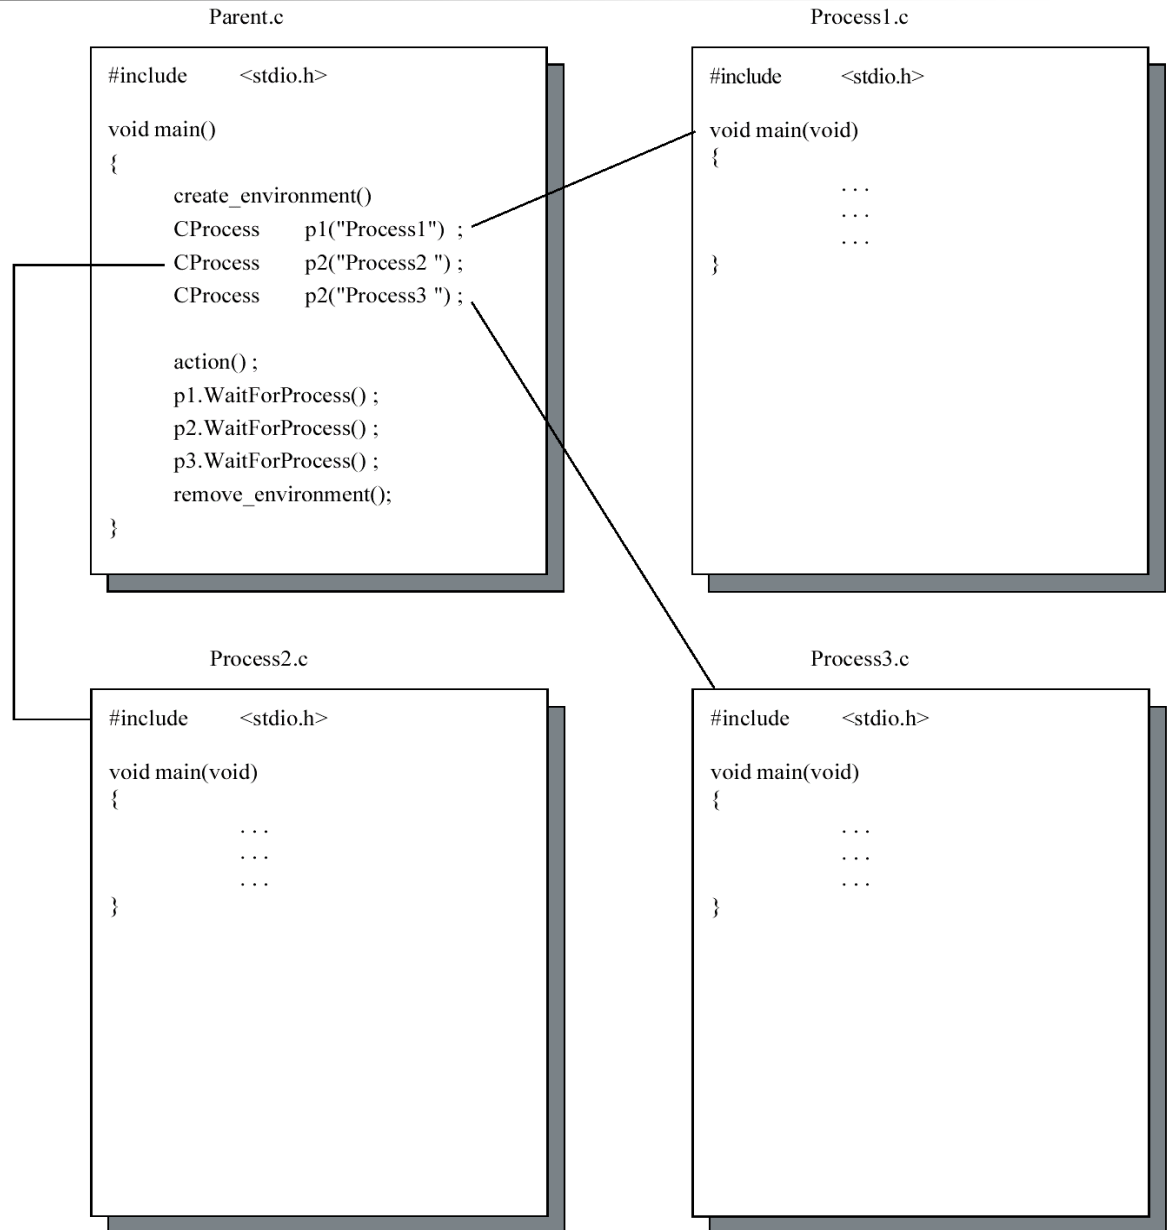
- At a fundamental, most crude level of parallel programming, we are interested in creating processes. A process is defined as any program that is currently **running** on the system, for example an editor, word-processor, compiler etc are all programs that could become processes if we ran them.
- In fact, the same program can be run **multiple** times on the systems creating multiple processes based on the same executable program.
- Creating processes with the **rt.cpp** library is pretty easy. One only has to create an **instance** of the **CProcess** class for a new process (i.e. a program) to be run for you on the host operating system. This is done via code in the **constructor** for the CProcess class.
- An example is shown below where a single '**Parent**' process attempts to create three '**child**' processes by creating three instances of the **CProcess** class, **p1**, **p2** and **p3**. The '**...**' is explained later.

```
void main()
{
    Create_Environment() ;           // create environment for child processes to run
    CProcess p1("Process1", ....) ; // create a child process
    CProcess p2("Process2 ", ....) ; // create a child process
    CProcess p3("Process3 ", ....) ; // create a child process

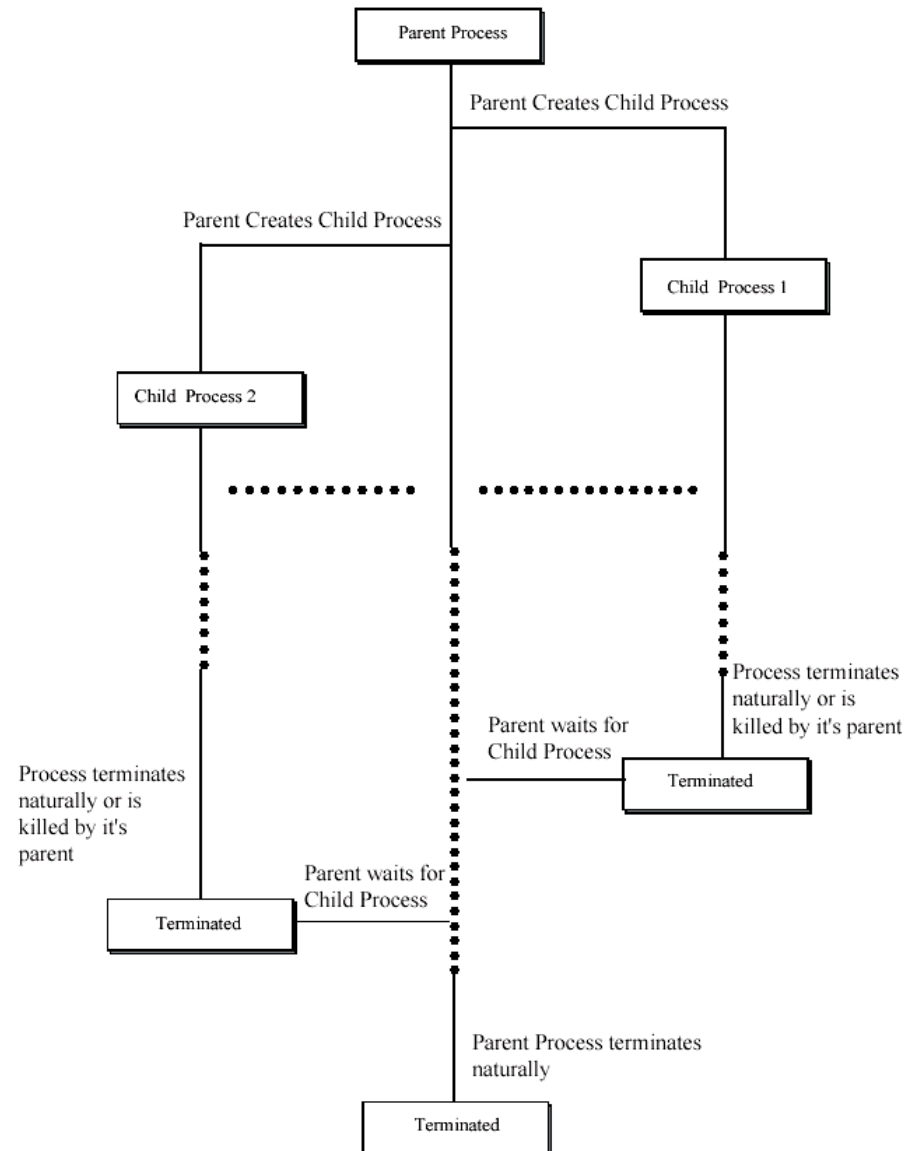
    action() ;                       // other actions carried out by this parent process

    p1.WaitForProcess() ;           // wait for child process p1 to terminate
    p2.WaitForProcess() ;           // wait for child process p2 to terminate
    p3.WaitForProcess() ;           // wait for child process p3 to terminate
    Delete_Environment() ;           // remove environment for child processes
}
```


- The illustration opposite relates the file name used in the **constructor** call for CProcess to the name of an **executable** program located somewhere on **disk**.
- Each **process** is thus represented by a **source file** with a single function **main()** which has been compiled to an executable program that can run, either on its **own**, or, under the control of a parent process.
- In effect, creating new CProcess objects, effectively asks the operating to locate the **' .exe' file** and start to run the function **main()** within it.
- The **environment** created by the parent represent any **resource** that the child processes expect to find in place when they begin execution, e.g. **files, process communication and synchronisation mechanisms** etc.



- The illustration opposite shows how with the creation of each new **CProcess object** a new operating system **thread** is created.
- This **thread** represents a 'line or trace of execution' that commences with the child programs function **main()**.
- All such threads are effectively **time sliced**, i.e. scheduled such that the CPU swaps rapidly between them so that each receives some CPU time and is seen to execute.
- When a process terminates, either **voluntarily** or when it is **terminated** by its **parent**, the thread of execution in that process is destroyed and this it is no longer scheduled to receive any CPU time.
- Waiting for a child causes the parent to suspend itself until any **one** of its child processes terminate. That is under Win32, you cannot specify that you wish to wait for process '**X**' to terminate.
- The best you can do is wait for any of them and then investigate afterwards which one terminated.



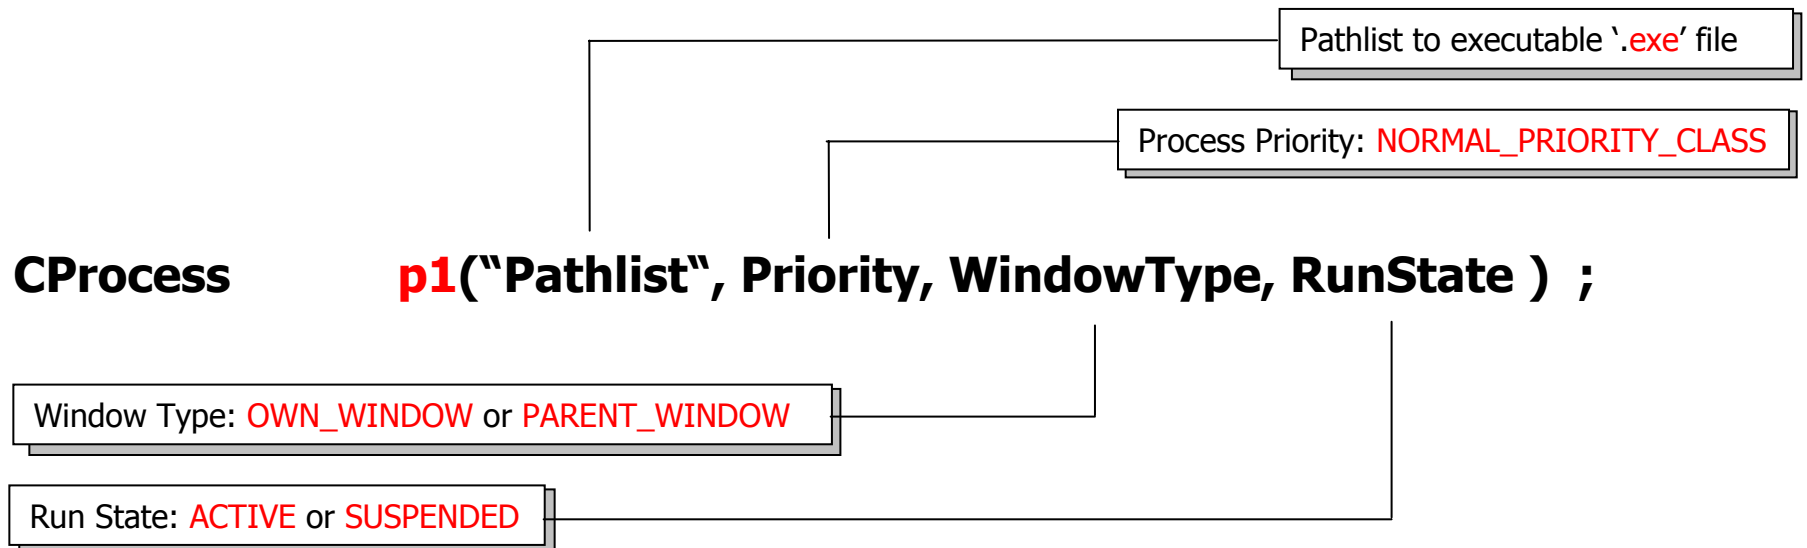
A More Detailed look at the CProcess Class

- The **CProcess** Class Encapsulates a number of **member functions** to facilitate the **creation** and **control** of a number of child processes.
- These member functions are outlined below with a brief description of what they do.
- A more detailed description and implementation of them can be found in the **rt.h** and **rt.cpp** files and the Tutorial Guide to Using Win32 (posted on the course web site)

CProcess()	-	The constructor responsible for creating the process
Suspend()	-	Suspends a child process effectively pausing it.
Resume()	-	Wakes up a suspended child process
SetPriority(int value)	-	Changes the priority of a child process to the value specified
Signal(int message)	-	Posts a message to a child processes (see later lecture)
TerminateProcess()	-	Terminates or Kills a child process (potentially dangerous)
WaitForChild()	-	Pauses the parent process until a child process terminates.

A More Detailed look at the CProcess Constructor

- The **CProcess** class constructor is responsible for
 - Locating the executable **.exe** program on disk via the specified pathlist.
 - Invoking the operating system kernel to ask it to run the program as a process.
 - Assigning it a **Priority** relative to all other processes in the system.
 - Assigning it a **Window** in which the programs I/O will interact.
 - Assigning it a run state: **Suspended** or **Running**.
- A detailed breakdown of this function call is given below. It takes **4** parameters



Example Detailed Program (see Q1 in tutorial)

```
#include "rt.h"

int main()
{
    CProcess p1("c:\\users\\paul\\parent\\debug\\paul1", // pathlist to '.exe.' file
                NORMAL_PRIORITY_CLASS,                  // a safe priority level
                OWN_WINDOW,                             // process uses its own window
                ACTIVE                                   // create process in running state
    );

    ...
    ...
    SLEEP(2000); // Pause parent for 2 seconds

    p1.Suspend(); // suspend the child process
    ...
    p1.Resume(); // resume the child process
    ...
    p1.WaitForProcess(); // pause parent until child terminates
};
```

How Does CProcess() Work?

- The code for the constructor for the **CProcess** class is given below, you can see how complex and tricky it is and how much detail it hides away from the programmer

```
CProcess::CProcess( const string &Name, int Priority, BOOL bUseNewWindow, BOOL bCreateSuspended):
    ProcessName(Name)
{
    STARTUPINFO StartupInfo = {
        sizeof(PROCESS_INFORMATION) ,
        NULL ,                      // reserved
        NULL ,                      // ignored in console applications
        (char *)(Name.c_str()) ,    // displayed in title bar for console applications
        0,0,                        // dwx, dwy, offset of top left of new window relative to top left of screen in pixel
        0,0,                        // flags below must specify STARTF_USEPOSITION. Ignored for console apps'
        0,0,                        // dwxsize, dwysize: Width and height of the window if new window specified
        0,0,                        // must use flags STARTF_USESIZE. Ignored for console apps'
        0,0,                        // size of console in characters, only if STARTF_USECOUNTCHARS flag specified,
        0,                          // Ignored for console apps
        0,                          // Colour control, for background and text. Ignored for console apps
        0,                          // Flags. Ignored for console applications
        0,                          // ignored unless showwindow flag set
        0,
        NULL,
        0,0,0                      // stdin, stdout and stderr handles (inherited from parent)
    };
    ...
}
```



Principles of Concurrent Systems: Processes & Threads

15

```
UINT flags = Priority ;           // Priority,

if(bUseNewWindow == OWN_WINDOW)  // if parent has specified child should have its own window
    flags |= CREATE_NEW_CONSOLE ;

if(bCreateSuspended == SUSPENDED) // if parent has specified child process should be suspended
    flags |= CREATE_SUSPENDED ;

BOOL Success = CreateProcess(    // CALL KERNEL HERE. This is where child process begins to run
    NULL,                        // application name
    (char *)(Name.c_str()),      // Command line to the process if you want to pass one to main() in
                                // the process
    NULL,                        // process attributes
    NULL,                        // thread attributes
    TRUE,                        // inherits handles of parent
    flags,                       // Priority and Window control flags,
    NULL,                        // use environment of parent
    NULL,                        // use same drive and directory as parent
    &StartupInfo ,              // controls appearance of process (see above)
    &pInfo                       // Stored process handle and ID into this object
);

ProcessHandle = pInfo.hProcess ;  // handle to the child Process, can be used to identify a process
ThreadHandle = pInfo.hThread;    // handle to the child process's main thread, used to identify thread

ProcessID = pInfo.dwProcessId ;  // Id of the Child Process (not the same as a handle)
ThreadID = pInfo.dwThreadId;     // Id of the Child Process's main thread (not the same as a handle)
}
```