

Parallel Processing and Process Granularity

- Programming a complex concurrent system making use only of multiple individual **processes**, represents what is known in concurrent programming terms as **crude** or **coarse grained process granularity**.
- That is, if we imagine our overall system as a **big rock** that we have to crack; what we have done is break down the rock into a number of **smaller rocks** so that each exists separately. These smaller rocks represent our **concurrent processes**.
- The problem is that in a big system, these smaller rocks could still be '**big**' (it's all relative)
- In theory we could just keep going, re-applying the principle of **hierarchical process decomposition** to break down our smaller **rocks/process** into even smaller ones so that there is even more concurrency in our system. In other words we are achieving even **finer grained process granularity**.
- Our ideal scenario would be to achieve '**dust**'. That is, we break down our rocks into smaller and smaller ones until they become **powder** and we cannot discern them individually any more. In process terms this would imply **perfect fine grained granularity**, where everything is running in parallel.
- In reality there is a **practical limit** imposed as to how **fine a granularity** we can achieve just by decomposing processes into ever smaller ones.
- Fundamentally this limit is a direct result of the **large overheads** introduced into a system in terms of **memory requirements** and **process swap time**.
- Furthermore, communication between processes becomes more difficult as an operating system does not allow one process to **share the same memory** used by another one, effectively building a **brick wall** around each process so that they cannot see or talk to each other.

Parallel Processing with Threads

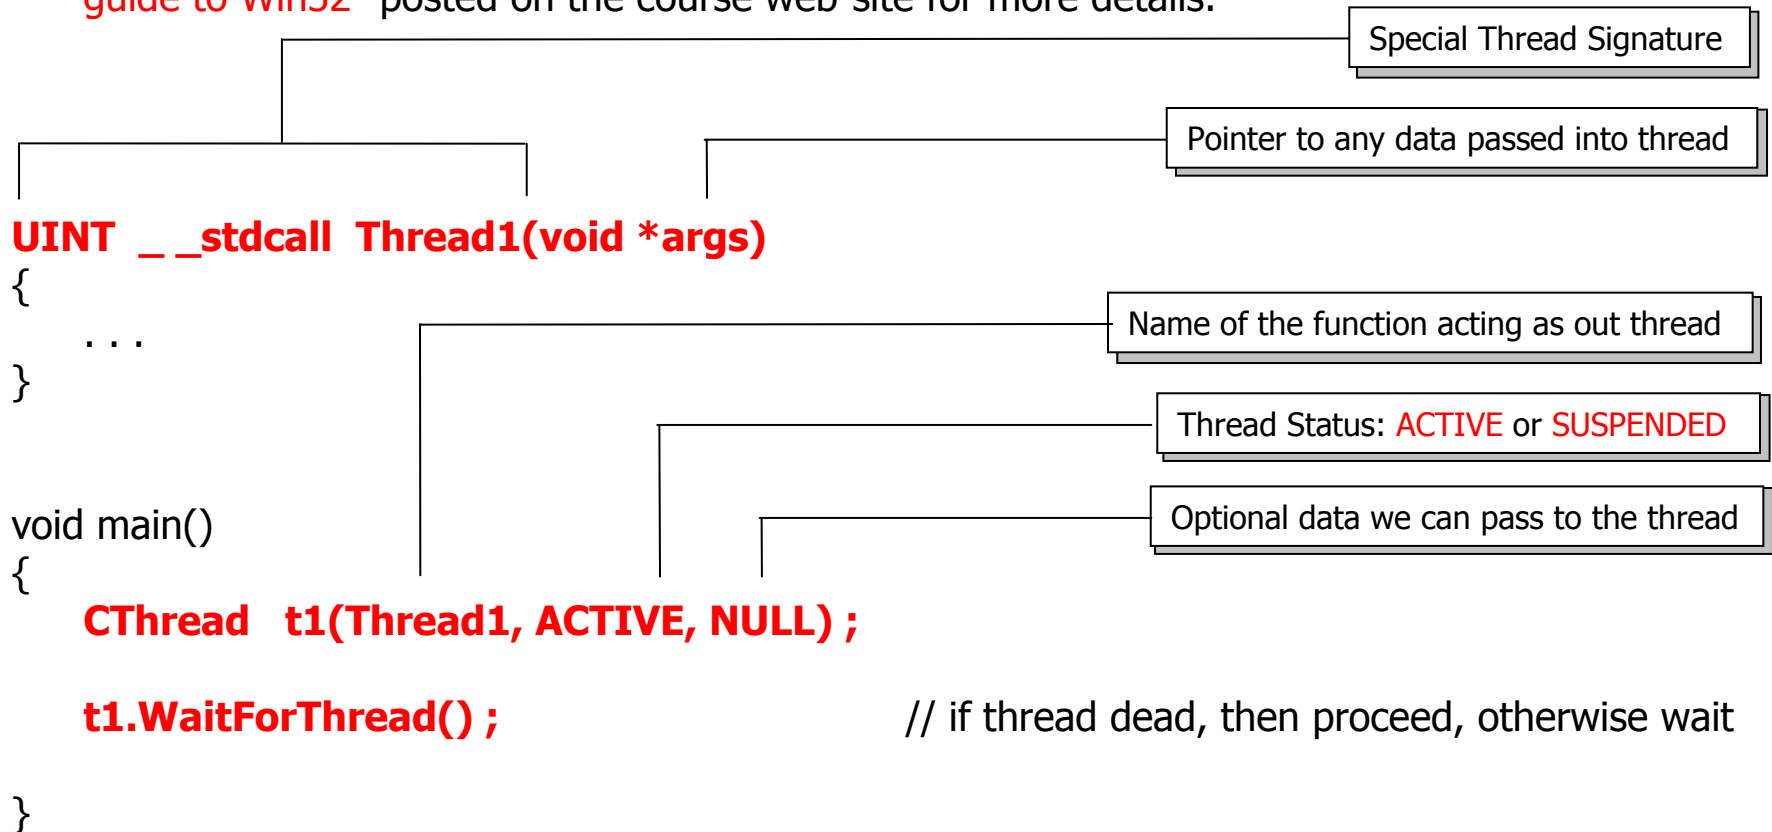
- An alternative, more practical approach is to introduce multiple '**threads**' into our processes in order to achieve the **desired fine grained concurrency**.
- Such threads impose less of an overhead as they are part of the same process (i.e. part of the same source code file) hence reducing **memory requirements** and the **time** to **swap** between them.
- In essence a **thread** is a parallel executing **section** of a programs code.
- It should be noted that every **process** has **at least one thread** that begins with the function **main()** when the program is run.
- From then on, a process can create as many other threads as it likes and these can all run concurrently (i.e. they are scheduled by the operating system and hence timesliced)
- Threads can be put to good use in breaking down activities within a process that should be executed in **parallel**. For example, imagine a simple data acquisition system. We could have threads to do the following
 1. A thread to monitor the external interfaces to the system e.g. 8 Analogue channels recording pressure, temperature etc, or perhaps 32 digital inputs.
 2. A thread to convert the above data into a response and generate an output.
 3. A thread to display graphically, to an operator, what is going on inside the system.
 4. A thread to archive the data to disk

Parallel Processing with Threads (cont...)

- Not all operating systems support the concept of thread level programming. The **Win32 Kernel** does as does **Linux**.
- In the right hands they can be a powerful tool. For example, Microsoft Office, a single application/process is multi-threaded, which means that for example **a file can be saved, printed and updated** on the screen all at the same time without you having to wait for each to finish.
- The Win32 kernel and Linux can even allocate individual **CPUs** to separate **threads** within a process if more than one CPU exists in the system.
- This is why some **Numerically Intensive** (i.e. they need lots of processing power) programs like **Adobe Photoshop** can show considerable speed improvements when running on a PC fitting with 2 CPUs.
- Unfortunately, a **single threaded** program running in a **multiple CPU** environment will show **little if any** improvement in speed and may even slow down.
- Hyper-Threading in Pentium Processors?

Creating Multi-Threaded Procedural Programs

- Each **thread** is written within the **same source file** for that process as if it were a simple **function** or **subroutine**, but it has a slightly modified **signature** (see below)
- Create a **CThread** object to represent our thread. The **constructor** for this class invokes the necessary calls to the kernel to create and schedule a thread running through the function.
- The example below, for a Win32 Application demonstrates the idea. See the handout "**A Tutorial guide to Win32**" posted on the course web-site for more details.



A More Detailed look at the CThread Class

- The **CThread** Class Encapsulates a number of **member functions** to facilitate the **creation** and **control** of a number of child threads.
- These member functions are outlined below with a brief description of what they do. They are similar to those of the CProcess class
- A more detailed description and implementation of them can be found in the **rt.h** and **rt.cpp** files.

CThread()	-	The constructor responsible for creating the thread
Suspend()	-	Suspends a child thread effectively pausing it.
Resume()	-	Wakes up a suspended child thread
SetPriority(int value)	-	Changes the priority of a child thread to the value specified
Signal(int message)	-	Sends a message to a child thread (see later lecture)
TerminateThread()	-	Terminates or Kills a child thread (potentially dangerous)
WaitForThread()	-	Pauses the parent process until a child thread terminates.

A More Detailed Example Program using Multiple Threads (See Q2 Tutorial for Example)

```
#include    "rt.h"

UINT  __stdcall Thread1(void *args)                // 1st thread function
{
    for(int i = 0; i < 1000; i ++)
        printf("Hello From Thread 1\n") ;

    return  0 ;
}

UINT  __stdcall Thread2(void *args)                // 2nd thread function
{
    for(int i = 0; i < 1000; i ++)
        printf("Hello From Thread 2\n") ;

    return  0 ;
}

void main()                                         // primary thread within this process
{
    CThread      t1(Thread1, ACTIVE, NULL) ;        // create two active secondary threads
    CThread      t2(Thread2, ACTIVE, NULL) ;

    t1.WaitForThread() ;                          // if thread already dead, then proceed, otherwise wait for it to finish
    t2.WaitForThread() ;                          // if thread already dead, then proceed, otherwise wait for it to finish
}
```

Creating Multiple Threads using one thread function: **PerThreadStorage**.

```
#include    "rt.h"
int    ThreadNum[8] = {0,1,2,3,4,5,6,7} ; // an array thread numbers, one for each thread

PerThreadStorage int MyThreadNumber ;    // How to allocate storage to each individual thread in the process including main thread

UINT __stdcall ThreadFn1(void *args)    // thread function
{
    MyThreadNumber = *(int *)(args);
    for(int i = 0; i < 100; i ++ )
        printf("Thread [%d]:. Monitoring Bit [%d]\n", MyThreadNumber, MyThreadNumber) ;

    return 0 ;
}

int main()
{
    CThread *Threads[8] ;

    // Now here is the clever bit with threads. Let's create 8 instances of the above thread code and let each thread know which port to monitor.
    for(int i = 0; i < 8; i ++ ) {
        printf("Parent Thread: Creating Child Thread 1 in Active State\n") ;
        Threads[i] = new CThread(ThreadFn1, ACTIVE, &ThreadNum[i]) ;
    }

    // wait for threads and then delete thread objects we created above
    for( i = 0; i < 8; i ++ )
        Threads[i]->WaitForThread() ;

    return 0 ;
}
```

Extract this threads number from its argument

An array of **CThread** Pointers

Create 8 new Threads giving each a number

Concurrent Programming in an Object Oriented World

- In object oriented languages like **Java** and **C++**, we can create **multiple threads** within our process through the more elegant use of **active objects**, i.e. objects with their own function **main()** (in C++) or function **run()** (in java) that have their own thread of execution running through them. Such objects execute concurrently with all other similarly '**active**' objects in the system.
- In **Visual C++** (using the rt.cpp library and the Win32 Kernel) we could derive our classes from the base **ActiveClass** as shown below. (See **Q2A** tutorial for an example)
- Next we override the virtual function **main()** inherited from that class to do whatever we want our class object to do. Finally we simply create **instances** of the class to create the threads.
- These threads are controlled as before. Note also that such threads are created in suspended state

```
class MyActiveClass : public ActiveClass {
    ...
private:

    // Must override main() inherited from ActiveClass. The base class constructor will then
    // create a thread to run though the function main()

    int main(void)
    {
        for(int i = 0; i < 1000; i++)
            printf("Say Hello to My Active Class.....\n") ;

        return 0 ;
    }
};
```




Principles of Concurrent Systems: Processes & Threads

24

```
class MyActiveClass1 : public ActiveClass {
private:
    int main(void) {                                // a thread within my c lass object
        for(int i = 0; i < 1000; i ++)
            printf("Say Hello to My Active Class 1.....\n") ;

        return 0 ;
    }
};

class MyActiveClass2 : public ActiveClass {
private:
    int main(void) {                                // a thread within my class object
        for(int i = 0; i < 1000; i ++)
            printf("Say Hello to My Active Class 2.....\n") ;

        return 0 ;
    }
};

int main(void)
{
    MyActiveClass1    object1 ;                    // create an instance of the above class
    MyActiveClass2    object2 ;                    // create an instance of the above class

    object1.Resume() ;                               // allow thread to run as it is initially suspended
    object2.Resume() ;                               // allow thread to run as it is initially suspended

    object1.WaitForThread() ;
    object2.WaitForThread() ;

    return 0 ;
}
```

Threads within a Class

```
class MyClassWithThreads : public ActiveClass {  
    ...  
    int PrintMessageThread(void *ThreadArgs) {  
        for(int i = 0; i < 10000; i++)  
            printf("%s\n", (char *)ThreadArgs);  
    }  
  
    int DisplayClassData(void *ThreadArgs) {  
        ...  
    }  
}  
  
int main(void) {  
    ClassThread<MyClassWithThreads> My1stThread( this, PrintMessageThread, ACTIVE, "Message 1" );  
    ClassThread<MyClassWithThreads> My2ndThread( this, PrintMessageThread, ACTIVE, "Message 2" );  
    ClassThread<MyClassWithThreads> My3rdThread(this, DisplayClassData, SUSPENDED, NULL);  
  
    // wait for the above active threads to complete  
  
    My1stThread.WaitForThread();  
    My2ndThread.WaitForThread();  
  
    // resume the 3rd thread and wait for it to complete  
    My3rdThread.Resume();  
    My3rdThread.WaitForThread();  
  
    return 0;  
}  
};  
  
int main()  
{  
    MyClassWithThreads Object1;  
    Object1.Resume();  
}
```