


Inter-Process Communication - The Pipeline

- One inter-process communication mechanism that handles synchronisation for itself is the **Pipeline**. This form of communications is generally characterised by the following properties:-
 - Implemented typically as a **wrap around buffer** with a **finite size**, i.e. a **circular queue**, which is used to store the data being communicated.
 - **Accesses** to the pipeline are handled via **software primitives**, i.e. **calls** to the **kernel** rather than the have the process **directly** access the data for itself.
 - This important difference means that access to the pipeline can be **managed** or **controlled** by the kernel that can prevent two or more processes accessing it at the same time.
 - Pipelines operate on a **FIFO/sequential** basis, i.e. data can only be read in the order it was written.
 - That is, **random and/or direct access** to data is not catered for, which makes it more amenable to implementation in a **distributed** system using a network. (The kernel accepts the read/write requests from the process and can transmit them serially over a network).
 - Restricted to communicating between **two processes** only, one reading, the other writing. Multiple pipelines will be required in any situation where the same data has to be transmitted to several reader processes.
 - The **read** operation is **destructive**, that is, once data has been **read** from the pipeline, it is physically **deleted** and no longer exists in the pipeline.
 - Theoretically data can only be **transferred in one direction**, from writer to reader, however most operating systems implement **bi-directional** pipelines by having two separate pipelines hidden within one.

Pipeline Synchronisation Properties

- Any process attempting to **write** to a **full Pipeline** will be **suspended** by the kernel until such time as there is space available to complete the write. This is to prevent data being lost if the circular buffer were to overflow and attempt to wrap around on itself
- Any process attempting to **read** data from an **empty buffer** will also be **suspended** by the kernel until there is sufficient information to **complete the transaction**. This is to prevent a reading process theoretically reading the **same data twice** when the buffer wrapped around.
- Finally, any process that has been **suspended** because it attempted to read from an empty buffer, or because it attempted to write to a full one will be allowed to **resume processing automatically** when the conditions that lead to its suspension have been removed i.e. the pipeline has data in it, or there is now space in the pipeline respectively.



Using a Pipeline

- Typically a pipeline is accessed using a number of **software primitives** with hooks to the operating system **kernel** to handle the **suspension** and **resumption** of processing. Such primitives might include the ability to:-
 - **Create** or **Open** a data pipeline.
 - **Read** data from a pipeline.
 - **Write** to a pipeline.
 - **Delete** a pipeline.
- The problem with **blindly reading** from or **writing** data to a pipeline is that a process runs the risk of getting **suspended** if the read or write cannot complete due to lack of **data** or **space** within the pipeline.
- This could be catastrophic in some systems, because a suspended process would mean that all other activities carried out by that process will also get suspended, not just the read or write operation. This could have a **ripple** effect throughout the system causing other process to suspend themselves and cause lockup in the system.
- For this reason, many operating systems provide a primitive to "**test the water**" w.r.t. the pipeline and see if a **subsequent** read or write operation would result in it being suspended. If necessary the process could then **defer** that read/write until later
- Thus many operating systems will also provide a primitive to :
 - **Determine** if there is **data** or **space** available to be read from or written to the pipeline.

A More Detailed look at the **CPipeline** Class

- The **CPipeline** Class encapsulates **five member functions** to facilitate the **creation** and **use** of a pipeline in a program.
- These functions are outlined below with a brief description of what they do. A more detailed description and implementation of them can be found in the **rt.h** and **rt.cpp** files.

CPipeline(Name, size)	-	The constructor responsible for creating the pipeline
BOOL Read(void *data, int size)	-	A function to read 'size' bytes of data from the pipeline and store at the address pointed to by 'data'. Returns true/false if the read is successful or if it fails. Note that suspension is not deemed a failure
BOOL Write(void *data, int size)	-	A function to write 'size' bytes of data to the pipeline using data stored in memory at the address pointed to by 'data'. Returns true/false if the read is successful or if it fails. Note that suspension is not deemed a failure
int TestForData()	-	Returns the number of available bytes of data in the pipeline that can be read without the process getting suspended
~CPipeline()	-	A destructor to delete the pipeline at the end of its use

- Tutorial Question **7** demonstrates the use of Pipelines

Example Use of Pipelines: Program 1 – Writing the Data

```
#include    "rt.h"

struct      example {                // structure template for data to be written to the pipeline
    int      x ;
    float    y ;
};

// Some data to be written in to the pipeline.

int  i = 5;                          // a simple int
int  array[ 10 ] = { 1 , 2 , 3 , 4 , 5 , 6 , 7 , 8 , 9 , 0 } ; // array of 10 integers

char name[15 ] = "Hello World" ;     // a string of 15 characters
struct example mystruct = {2, 5.5 } ; // a structure with an int and a float

int  main()
{
    CPipe  p1("MyPipe") ;           // Create a pipe 'p1' with the name "MyPipe"

    p1.Write(&i, sizeof(i)) ;        // write the int 'i' to the pipe
    p1.Write(&array[0], sizeof(array)) ; // write the array of integers' to the pipe
    p1.Write(&name[0], sizeof(name)) ; // write the string to the pipe
    p1.Write(&mystruct, sizeof(mystruct)) ; // write the structure to the pipeline

    return 0 ;
}
```

Example Use of Pipelines: Program 2 – Reading the Data

```
#include      "rt.h"
struct        example {                // structure template for data that we intend to read from pipeline
    int       x ;
    float     y ;
} ;

// Some variables to hold the read from the pipeline.

int   i ;                // a simple int
int   array[ 10 ] ;      // array of 10 integers

char  name[15 ] ;        // a string of 15 characters
struct example mystruct ; // a structure with an int and a float in it

int   main()
{
    CPipe   p1("MyPipe") ;           // Create a pipe 'p1' with the name "MyPipe"

    p1.Read(&i, sizeof(i)) ;          // Read the int 'i' from the pipe
    p1.Read(&array[0], sizeof(array)) ; // Read the array of integers' from the pipe
    p1.Read(&name[0], sizeof(name)) ; // Read the string from the pipe
    p1.Read(&mystruct, sizeof(mystruct)) ; // Read the structure from the pipeline


    // Now print out the data read from the pipeline

    printf(" i = %d\n", i) ;
    for( int x = 0; x < 10; x ++ )
        printf("array[%d] = %d\n", x, array[x]) ;

    printf("%s", name) ;
    printf("mystruct.x = %d, mystruct.y = %f\n", mystruct.x, mystruct.y) ;
    return 0 ;
}
```

Handling Multiple Pipelines

- Suppose a process is actively reading data from say **three** pipelines and maybe a keyboard (which has properties similar to a pipeline) as shown below. Because data could arrive along **any pipeline** at **any time**, the process will probably not be able to predict in advance which pipeline to read first.
 - If the process were to **make a guess** at which pipeline **holds data** by attempting to **read** from it, we know from the properties of a pipeline that it will be suspended if there is **insufficient data** to satisfy the read operation resulting in the process being unable to read data from **any other pipeline** when that arrives.
 - These pipelines would all eventually **fill up** and lead to their **writers** also getting suspended. How do we solve this ?
1. Use the **TestForData()** Primitive
 2. Use **Multiple Threads**, one for each pipeline



Solution using TestForData()

```
#include      "rt.h"

int   pipe1data ;                // a single integer
float pipe2data[10] ;            // an array of 10 floats
char  pipe3data[20] ;            // a string of up to 20 character
char  KeyData[2] = {'\0', '\0'} ; // a 2 character keyboard command initialised to be empty

int   main()
{
    CPipe   p1("Pipe1") ;        // create the three named pipelines
    CPipe   p2("Pipe2") ;
    CPipe   p3("Pipe3") ;

    // Now generate an endless polling loop checking if any data has arrived from any of the four sources. No error checking for the sake of clarity.

    while( 1 )
    {
        if ((p1.TestForData()) >= sizeof(pipe1data)) // if at least 1 integer in pipeline
            p1.Read(&pipe1data , sizeof(pipe1data)) ; // read data from pipe

        if ((p2.TestForData()) >= sizeof(pipe2data)) // if at least 10 floats in pipeline
            p2.Read(&pipe2data , sizeof(pipe2data)) ; // read data from pipe

        if ((p3.TestForData()) >= sizeof(pipe3data)) // if a 20 character string in pipeline
            p3.Read(&pipe3data , sizeof(pipe3data)) ; // read data from pipe

        // The primitive TEST_FOR_KEYBOARD() below tests the keyboard to see if any key
        // has been pressed, if so it returns true (i.e. a value other than zero)

        if (TEST_FOR_KEYBOARD() != 0) { // if a key has been pressed maintain a scrolling array of the last two characters read
            KeyData[0] = KeyData[1] ; // move up previous character read
            KeyData[1] = getch() ;    // read the character from keyboard
        }

    }
}
```


Type Safe Pipelines

- The successful operation of a pipeline depends very much on the reader and writer process reading and writing **the same type of data in the same order**, e.g. an 'int', followed by a 'float', then a 'double' etc.
- For example, it would be useless and cause chaos in the system if the writer process wrote a 'float' into the pipeline, but the reader process read that data as an 'int', or perhaps 4 'chars'.
- As far as the pipeline is concerned, there is no difference between the data, since they are both (typically) 4 bytes in size and as long as the writer writes a 4 byte item of data into the pipe, and the reader reads the same sized data out, then it's not the pipelines problem !!
- To avoid these sorts of problems we could introduce an element of **type-safety** into our programs, in essence we could use a more sophisticated version of the pipeline which only permits data to be read/written according to a **template** agreed when the **pipeline is created**.
- For example, the template would state that this pipeline can only be used to hold integers, or float or objects/structures or type 'X'
- Such a pipeline is represented by a **C++ templated class version** of the Pipeline, called, not surprisingly **CTypedPipe**.
- An example of its use is shown overleaf

Example Usage of a Type safe Pipeline

```
#include "rt.h"

CTypedPipe<int>      Pipe1("pipe") ;           // create an 'int' pipe

int main()
{
    int x, i=5, j=0 ;

    for(x=0; x < 10; x++)      {
        Pipe1.Write(&x) ;           // note no size argument req'd as size is implied by type of pipeline
        printf("Wrote %d into Pipe. Number of ints in pipe = %d\n", x, Pipe1.TestForData()) ;
    }

    for(x=0; x < 10; x++)      {
        Pipe1.Read(&j) ;           // again no size argument req'd as size is implied by type of pipeline
        printf("Read %d from Pipe. Number of ints in pipe = %d\n", j, Pipe1.TestForData()) ;
    }

    return 0 ;
}
```

The type of data held by the pipeline