


### The General Solution to Solving Mutual Exclusion

- In attempting to come up with a solution to mutual exclusion in time-sliced (and multiple CPU) systems, we must isolate and analyse the fundamental **problem** that is causing the system to fail.
- In using **flags** to mark a resource as free or busy, the problem lies with the fact that a process can be **interrupted** part way through the normal testing and setting of a flag.
- If we are going to solve the problem, then we must ensure that the action of setting the flag is made **indivisible** or, to use the terminology of the subject area, made **'atomic'**, (that is, we cannot 'split' it any further) whilst at the same time ensuring that we implement a queuing mechanism to prevent indefinite lockout and thus provide deterministic behaviour).
- If we are to achieve this, then we must move the **flag** out of the control of each process and place it under the control of the OS. kernel itself, which is controlling the time slicing/multi-processing actions.
- The Kernel then provides primitives to allow the processes to access the flags **indirectly**.
- If calls to manipulate the flag are done via the kernel, it can **suspend time slicing** to ensure that the process cannot be interrupted during the testing and setting of the flag and at the same time **suspend a process** and place it in a **first-in, first-out queue** if the resource is suspended.

## The General Solution to Mutual Exclusion - The Binary Semaphore

- The Dutch mathematician **Dijkstra** is generally credited with solving the problem of mutual exclusion in computer systems.
- It is believed that he came up with the solution after studying a real life mutual exclusion problem; how trains share a **single track railway line**; a classic non-sharable resource if ever there was one.
- He observed that railway tracks were **partitioned** or split into **multiple sections**, so that many trains could actually share the track at once, but only one was allowed to use a particular critical section of that track. That is mutual exclusion need only be enforced on each section, not the whole track.
- Dijkstra observed that a Railway **Signal** was used to indicate when a section of single track was **free** or **busy**. A train, arriving at the signal would have to **wait** if the section was in use by another train already travelling along it (indicated by a **raised** signal).
- When a train left a particular section of track, it would **lower** the signal at the **start** of the section, thereby allowing the following train to enter that vacated section, which would immediately raise the signal again for the next train behind that.
- This prevented two or more trains using the same non-sharable section of a single track line.
- Dijkstras solution was to implement this Signalling concept within the Operating System **Kernel** using a **Binary Semaphore**.



### Properties of the Binary Semaphore

- It can be in one of two states: **Signalled** or **Not Signalled/Blocked** (0 and 1 respectively)
- It is a protected variable not **directly** accessible to any process.
- Two primitives **Signal()** and **Wait()** are available to a process that can be used to **indirectly** manipulate the binary semaphore.
- Of fundamental importance is the fact that **once started** these two primitives are **guaranteed to complete** without interruption or time slicing. That is they are **atomic**, or **indivisible** primitives.
- In a multi-processor system, this means that the operating system would implement these primitives using the **TAS** instruction, or, in a time sliced system, by suspending the **RTC** and using a **flag**.
- Furthermore, the OS Kernel can **queue** the processes wishing to use the resource and release them on a first come first served basis when the resource is released.
- In the **rt.cpp** the concept of a semaphore is encapsulated by the **CSemaphore** class.

### Operation

- To solve mutual exclusion, each process must be written to perform a **Wait()** before accessing the resource, followed by a **Signal()** after it has finished, thus each process looks like this.

#### **WAIT(S)**

access or update non-sharable data pool

#### **SIGNAL(S)**

### Implementing the **WAIT** Primitive

- A pseudo code implementation for the **Wait()** primitive is shown below.
- If a process executes a **Wait()** primitive on a semaphore **S** with the value **1**, then the process will be allowed to continue and the semaphore is then decremented by one.
- Otherwise the process is suspended until a Signal is performed on the semaphore later by another process.

### Pseudo code for **WAIT** primitive on semaphore (**S**)

```
If( S == 0 )  
    suspend process  
else  
    S = S - 1
```

### Implementing the **SIGNAL** Primitive

- A pseudo code implementation for the **Signal()** primitive is shown below.
- If a process executes a **Signal()** primitive on a semaphore, then if there are processes **waiting** to use the resource, then the process at the **head** of the queue is allowed to wake up from its suspended '**Wait()**' call and resume processing. In this case the value of the semaphore does not change
- However, if there are no processes waiting to use the resource, then the semaphore is reset back to 1

### Pseudo-code for the **SIGNAL** primitive on semaphore (**S**)

**If( there are any processes waiting on this semaphore )**

**let the first in the queue resume processing.**

**else**

**$S = S + 1$**



## Process and Thread Synchronisation Concepts

20

```
/******
**   Program 1 - Example use of a CSemaphore object to enforce mutual exclusion
**   *****/
#include      "rt.h"
struct      mydatapooldata {                // data to be stored in a datapool
    int x ;
    int y ;
};

struct mydatapooldata *DataPtr ;           // pointer to the datapool

int  main()
{
    // Start off with a datapool to represent our non-sharable resource.

    CDataPool  dp1("MyDataPoolName", sizeof(struct mydatapooldata)) ;
    DataPtr = (struct mydatapooldata *)dp1.LinkDataPool() ;

    // Create a semaphore 'mutex1' with initial value 1 meaning the resource is free
    CSemaphore s1("mutex1", 1) ;

    // Create a new process that runs concurrently with this one
    CProcess      p1("c:\\users\\paul\\parent\\debug\\program2") ;

    for( int i = 0; i < 10000; i ++ )        // access resource 10000 times
    {
        s1.Wait() ;                        // gain access to resource
        .....                               // code to access a non-sharable resource
        s1.Signal() ;                       // release resource
    }

    p1.WaitForProcess() ;                    // wait for child process to finish
    return 0 ;
}
```



## Process and Thread Synchronisation Concepts

21

```

/*****
**   Program 2 - Example use of a CSemaphore object to enforce mutual exclusion
*****/
#include      "rt.h"

struct          mydatapooldata{          // data to be stored in a datapool
    int          x ;
    int          y ;
};

struct          mydatapooldata   *DataPtr ;          // pointer to the datapool

int   main()
{
    //   A datapool to represent our non-sharable resource. Now we need to make sure

    CDataPool   dp1("MyDataPoolName", sizeof(struct mydatapooldata)) ;

    DataPtr = (struct mydatapooldata *)dp1.LinkDataPool() ;

    //   Create a semaphore 'mutex1' with initial value 1 meaning resource is free

    CSemaphore s1("mutex1", 1) ;

    for( int i = 0; i < 10000; i ++ )          // access resource 10000 times
    {
        s1.Wait() ;          // gain access to resource
        .....          // code to access a non-sharable resource
        s1.Signal() ;          // release resource
    }

    return 0 ;
}

```

## Using a **Mutex** as an Alternative to a **Semaphore**

- Win32 also supports the concept of a mutex (encapsulated by the **CMutex** object in **rt.cpp** Library)
- Essentially, this is a **thinly disguised semaphore** which is always created with a value of **1**.
- More subtle is the fact that a mutex can **only** be **signalled** by the process that has acquired the resource, that is, the one that performed the **Wait()**.
- In addition a process that performs a **Wait()** on a mutex protecting a resource that it has already acquired will not block itself but it has to remember to do the corresponding number of **Signal()** operations to release it afterwards.
- This is to prevent another process from **illegally unlocking** a resource by signalling it when it does not even own that resource (more likely to happen to by accident due to a bug than by deliberate intent)

```
#include "rt.h"

int main()
{
    ...
    CMutex m1("mutex1"); // create a mutex
    ...

    for( int i = 0; i < 10000; i ++ ) // access resource 10000 times
    {
        m1.Wait(); // gain access to resource
        ... // code to access a non-sharable resource
        m1.Signal(); // release resource
    }
    ...
}
```



### Mutual Exclusion between Threads: **Critical Sections**

- Although a semaphore or mutex are guaranteed to successfully implement mutual exclusion between two competing processes (or indeed between two threads) the use of either introduces considerable overheads affecting the execution speed of each process.
- In effect the use of a Mutex or a semaphore makes the programs run **much slower** than if semaphores had not been used.
- In a situation where we are attempting to enforce mutual exclusion between **two threads**, within the SAME process, we can use a much faster method based around the concept of a '**critical section**' (supported by a CriticalSection object in **rt.cpp** library).
- The example on the next page demonstrates the concept.



# Process and Thread Synchronisation Concepts

24

```
#include      "rt.h"

//  A GLOBAL critical section object MUST can be created that can be accessed
//  by all threads within the same process

CriticalSection cs ;           // global object created at start of process and destroyed at end

int  x ;                       // data shared between two threads, in this case a simple 'int'
```

```
UINT  __stdcall Thread1(void *args)    // A child thread to update the datapool
{
    for(int i=0; i < 5000000; i ++){
        cs.Enter();                    // gain access to the resource via the critical section
        x = x + 1 ;                    // update x
        cs.Leave();                   // release the resource
    }
    return 0 ;
}
```

```
//  The main parent thread in the process
```

```
int  main()
{
    CThread  Child( Thread1 ) ;        // create a child thread

    for (int i=0; i < 5000000; i ++){
        cs.Enter();                    // gain access to the resource via the critical section
        x = x + 1 ;                    // update x
        cs.Leave();                   // release the resource
    }

    Child.WaitForThread();              // wait for child thread to finish
    return 0 ;
}
```

### Monitors

- The problem with using a mutex/semaphore to enforce mutual exclusion is that its success relies on all processes and threads being written correctly to use the mutex/semaphore before accessing the resource and releasing it afterwards.
- Any **bugs** or **bad behaviour** on the part of the process (such as forgetting to release the resource or performing a **Wait()** twice without releasing it) could mean that the whole system becomes locked with every process blocked and nobody able to access the resource.
- The solution proposed by 'C.A.R. Hoare' was the introduction of a **Monitor**, which is a collection of **functions/subroutines** and **data** all packaged into a special kind of **module** whose data is **hidden** from **direct access** by a process/thread. A process can only gain access to that data through the use of an **interface function** provided by the monitor.
- If this sounds like a Java/C++ class to you then you are thinking along the right lines, but in 1975 when it was first proposed, object oriented programming did not exist and the concept of **information hiding** and **data/function encapsulation** was a new concept then.
- The **interface functions** could then be written to perform the **Wait()** and **Signal()** operations required to gain access to the monitor data automatically and transparently whenever a process/thread entered the monitor and left it.
- This removed the responsibility from the process/thread for implementing the correct protocol to gain access to the data, i.e. the **Wait()** and **Signal()**, and delegated it to the monitor and as a result access to the data became safer.

### An Example C++ code Monitor to implement **Mutual Exclusion**

```
class      MyDataMonitor {
private:
    int value ;                // resource to be protected, i.e. an integer in this case, but it could be anything
    CMutex m1(....) ;         // the mutex to protect the data

public:
    int read()
    {
        m1.Wait()             // gain sole access to the data
        int x = value ;       // access the data
        m1.Signal() ;         // release access to the data
        return data ;
    }

    void write(int TheNewData)
    {
        m1.Wait()             // gain sole access to the data
        value = TheNewData ;   // update the data
        m1.Signal() ;         // release access to the data
    }
}
```

Placing a monitor object inside a **datapool** allowed multiple processes to share it.

### Java Monitors

- Java takes the concept of a Monitor one stage further by hiding the existence of the mutex behind the concept of a **synchronized method**, which **implies** the existence of a mutex.
- Execution of a synchronized method by a thread means that **Wait()** and **Signal()** operations will be automatically on the hidden mutex.
- This approach ensures that only one thread is ever permitted inside **any synchronized method** for the class at any one point in time.
- If the **whole class** is declared **synchronized** then all its methods are synchronized methods

```
public class MyDataMonitor {  
    private int value ;           //data to be protected  
  
    public synchronized int read()  
    {  
        return data ;  
    }  
  
    public synchronized write(int TheNewData)  
    {  
        value = TheNewData ;    // update the data  
    }  
}
```