

Strategy 4 - Deadlock Prevention – The Banker's Algorithm

- The main algorithms for preventing deadlock are based upon the concept of *safe and unsafe states*.
- Such algorithms avoid deadlock by *preventing a circular list of dependencies forming* between a set of processes and the resources they are attempting to acquire. (Remember that this was one of the 4 '*Coffman*' conditions required for a system to be in a deadlock state).
- Preventing this condition arising thus prevents the possibility of deadlock occurring. Stated simply, you deny a process's request to acquire a resource if it would lead to a circular dependency list forming in the first case.
- In essence the system always tries to maintain itself in a *safe state* where by there are always enough *free resources* to satisfy *at least one process*, i.e. there is always at least one process that can *proceed* and is not *blocked*.
- Dijkstra came up with a scheduling algorithm/solution known as the *Banker's algorithm* because it was modeled upon the way that a small town banker might deal with a group of customers to whom he has granted *credit or overdraft arrangement*. Such a scheme can be employed by any operating system wishing to prevent deadlock occurring. Let's see how it works.

Strategy 4 - The Banker's Algorithm – an Explanation

- Imagine that a **banker** (in this analogy our **operating system**) has several **customers** (i.e. **processes**) which need to borrow certain sums of **money** (i.e. **resources**) which they intend to repay (i.e. they will return the resource)
- For example, let's suppose the banker has **4** customers, we'll call them **A, B, C** and **D**, each of which have agreed, **in advance** with the banker, that they can borrow a certain amount of money in the future i.e. they have agreed in principle to a **loan** or **overdraft**.

(This is analogous to our processes negotiating in advance with the operating system that they intend to borrow several resources in the future).

- Let's suppose that the banker has agreed in principle that
 - **A** can borrow \$**6000**
 - **B** can borrow \$**5000**
 - **C** can borrow \$**4000** and
 - **D** can borrow the most at \$**7000**,
- That is, the banker has agreed in principle to loans totaling \$**22000**

Strategy 4 - The Banker's Algorithm – an Explanation (cont...)

- The problem here (which is an important point in the discussion on deadlock prevention) is the fact that the banker **doesn't actually have \$22000** to lend. At least not all at once !!.
- In fact he only has \$**10000**, but this is still **more than sufficient** than the **maximum** amount he has agreed to lend to any **one** of his customer (Customer **D** at \$**7000**)
- In effect the banker is **gambling/relying** on the fact that not **all** of the customers will want to borrow **all** their entitlement **all** at the **same time**.
- In other words he is betting that the loans will be **phased** at different times and that some of them may actually pay back **some** or **all** of their **loans** before he runs out of money to lend.
- In our computing analogy, this is equivalent to saying that an operating system has **10 resources**, but that processes **A, B, C** and **D** might actually want **22** if they all request those resources simultaneously. In other words there is the potential for a deadlock, which may or may not happen depending upon the order and frequency with which the banker lends and received money.
- However the banker and the operating system still have enough money/resources in reserve at all times to meet each individual customers/processes request for its maximum remaining allocation of money/resources.
- Let's see what happens

Strategy 4 - The Banker's Algorithm – an Explanation (cont...)

- When the banker opens his bank on Monday morning, none of the customers have borrowed any money.
- This initial state of play is shown in Fig (a) below and at this point at least, the banker can satisfy **any single request** made **by any one customer** for their **full loan entitlement**.
- However, some of his customers may arrive at various times during the day and request some or perhaps their entire agreed loan (this is equivalent to several processes requesting some or all of their agreed resources from the operating system).
- Let's move on and consider Fig (b). Here, at say 3.30pm the banker has given \$1000 to Customers A and B, \$2000 to C and \$4000 to D. Thus he has loaned an amount totaling \$8000 and now has just \$2000 in reserve.
- The important point to note here is that there is still enough in reserve to satisfy **at least one** of his customers if they request the remainder of their full entitlement; in this case customer C who can borrow another \$2000.
- This state of affairs is referred to as a **safe state**, because at least one of his customers can proceed with their full loan i.e. they cannot be deadlocked because the banker has enough resources to allow at least one of his customers to borrow money.

Has Max		
A	0	6
B	0	5
C	0	4
D	0	7

Free: 10

(a)

Has Max		
A	1	6
B	1	5
C	2	4
D	4	7

Free: 2

(b)

Has Max		
A	1	6
B	2	5
C	2	4
D	4	7

Free: 1

(c)

Strategy 4 - The Banker's Algorithm – an Explanation (cont...)

- Let's suppose now that at 3.40pm, customer **B**, requests a further \$1000 of his agreed loan, as shown in Fig (c). This would mean that the banker now only has \$1000 in reserve.
- This is an **unsafe state**, because the banker does not have enough in reserve to satisfy any one of his customers if they call upon their full loan entitlement and a deadlock situation **could** arise.
- However, it does not mean that a deadlock will arise, it's just that the **potential exists** and to avoid it, the banker would temporarily refuse the loan of the £1k to customer **B** in Fig (c) until some other customer had paid back sufficient of their loan such that granting **B's** request would still leave the system in a **safe state**.
- The idea then is that a safe state is one where an operating system has sufficient resources in reserve to satisfy the full amount of outstanding resource requests that could be made by at least one process.
- If a request is made that would lead the operating system into an **unsafe state**, then that request is temporarily refused and the process will be put on hold.

Has Max		
A	0	6
B	0	5
C	0	4
D	0	7

Free: 10

(a)

Has Max		
A	1	6
B	1	5
C	2	4
D	4	7

Free: 2

(b)

Has Max		
A	1	6
B	2	5
C	2	4
D	4	7

Free: 1

(c)


Limitations of the Banker's Algorithm

- The main problem or limitation of this algorithm, as you have probably identified yourself, is the fact that processes have to declare **in advance** the **number** of **each type of resource** they are intending to use.
- If they cannot do this, then deadlock prevention using this scheme cannot be assured.
- Such a scheme works well in batch programming, but in interactive or real-time systems, where processes are created in response to users logging in or events triggering the creation of new threads (which may request resources) it cannot be employed safely and the risk of deadlock will continue.

Classical Deadlock Problems –


The Dinning Philosophers

- The picture opposite illustrates a classical **deadlock** and **starvation** problem known commonly as the Dinning Philosophers problem and was first proposed by Dijkstra as a vehicle for operating systems designers to test their deadlock detection and recovery scheme against
- It also serve a dual purpose in that it gives students of **operating systems** and **concurrent programming** something of a problem to keep them awake at night in their attempts to understand its solution.




The Dining Philosophers – An Explanation

- The original problem concerned **five philosophers** who sit around a circular table. Each philosopher has their **own bowl** and access to **two chopsticks** which are also shared with each of the philosopher's neighbours. In fact the chopsticks are the classical **resources** that each philosopher fights over.
- The life of a philosopher consists of alternate bouts of **eating**, for which he needs **two chopsticks**, and **thinking** which requires **none**. The problem here is one of both **deadlock** and **starvation**.
- Imagine all philosophers decide to **eat at the same time**, if they all pick up their **left** chopstick first, followed by the **right**, then deadlock will occur, ditto if they pick of their **right** first followed by their **left**.
- We could modify the approach so that a philosopher picks up their left chopstick and then checks if the right is available. If it is not, he puts down his left chopstick and tries again later.
- Unfortunately, this solution can lead to **starvation** as we saw with voluntary relinquishment of a resource, whereby if a philosophers gives up a chopstick he/she may never get it back again.




The Dinning Philosophers – A Solution

- The solution always presented in books is rather **terse** and **cryptic** and even though it is short, it is quite **difficult** to follow.
- An example of it, implemented as a class can be found in the **rt.cpp** library header file.
- Look for the class **CDinningPhilosophers**, while a working solution/implementation of it can be found in **Q12** of the tutorial questions.
- Even then, the solution is not very good and is actually quite **unrealistic** since it doesn't really involve philosophers attempting to acquire the resources and if they are busy releasing the ones they already possess (as they would in the real world), but rather involves a philosopher **testing the state of his neighbouring philosopher** to see whether they are eating or not and hence **deducing** whether or not the chopsticks are free, which isn't really the point.
- It's a bit like process A checking what process B is doing and hence deducing whether B is using the resource that A wants.
- Thus the whole solution isn't even a generic solution since it is a specific solution to a specific problem, i.e. the dinning philosopher problem
- A more **generic solution** is actually a lot simpler and can easily be tailored to solving the dinning philosophers problem as well as other problems and makes use of a specific **Win32 call** (which other operating systems will have in some form or another) called the **WAIT_FOR_MULTIPLE_OBJECTS()**



The Dinning Philosophers – A Solution (cont...)

- The **WAIT_FOR_MULTIPLE_OBJECTS** system call in the Win32 kernel simply allows a process/thread to request multiple resources simultaneously.
- The kernel will ensure that that none of the resources that the process/thread is requesting will be given to it until it can **simultaneously acquire them all**.
- In effect this behaves much like the deadlock prevention scheme discussed earlier, where multiple resources were treated as one resource. Either a process/thread has them all, or it has none of them.
- However this scheme has the added flexibility that other processes/thread can request a sub-set of those resources and thus does not suffer from the limitations of that scheme.
- It does not however eliminate the possibility of **starvation**, since even if the call manages to locate a subset of the requested resource as being available, it cannot acquire them until they are all available and thus the possibility exists that another process/thread will acquire the available ones first.
- If this happens frequently enough, then a process/thread may never find all its requested resource available and thus **'starve'**



The Dinning Philosophers – A Solution

- To demonstrate the use of the **WAIT_FOR_MULTIPLE_OBJECTS** call, imagine we have two separate resources, a scanner and a CD Writer with their corresponding two semaphores to protect them, as shown below.
 - CMutex Scanner(.....) ;
 - CMutex CDWriter(.....) ;
- Now imagine a process wants to acquire **both** of these resources without the possibility of **deadlock**. It obviously has to wait for both of them to become free and thus we cannot acquire one then the other sequentially, we have to acquire them both simultaneously. We could achieve this using the **WAIT_FOR_MULTIPLE_OBJECTS** call as shown below.
- First we have to create an array of '**HANDLES**' (these are for all intents and purposes identifies for the **mutex's**) to the resources we want to acquire; in this case 2.

```
HANDLE MyResourceHandles[2] = {  
    ScannerMutex.GetHandle(),           // get handle for each mutex  
    CDWriterMutex.GetHandle()  
};
```

The Dinning Philosophers – A Solution

- Having got the two handles to the two resources we could now attempt to acquire them thus

```
WAIT_FOR_MULTIPLE_OBJECTS(  
    2,                // number of resources to acquire  
    MyResourceHandles, // array of handles to those resources  
    INFINITE           // try for this length of time to acquire them or give up  
);
```

- Here, we give the function the array of handles to the resources we require and then tell it to acquire them for us. We have specified a time period of **INFINITE** meaning do not return until you have got them.
- Once the function returns, we know that we have sole access to the two resources and that a Wait() operation will have been performed on the resources so that nobody else can acquire them.
- We still of course have to Signal() the resources in the usual manner when we have finished with them.
- Armed with this, we could thus solve the dinning philosophers problem as shown below



Deadlock and Starvation

31

```
#define      NUMBER_OF_PHILOS      5                                // start off with 5 philosophers
CMutex      *Chopsticks[NUMBER_OF_PHILOS] ;                       // create an array of mutex's, one for each philosopher

// This function returns the index of the chopstick mutex to the left of the current philosopher
int LeftChopstick(int i)      { return ((i + NUMBER_OF_PHILOS - 1) % NUMBER_OF_PHILOS) ; }

// This function returns the index of the chopstick mutex to the left of the current philosopher
int RightChopstick(int i)     {return ((i + 1) % NUMBER_OF_PHILOS) ; }

// A thread to represent a philosopher
UINT _stdcall Philosopher(void *args)
{
    int      MySeatNumber = *(int *)(args) ;                        // figure out which philosopher I am

    // Now get the indexes into the array for my left and right chopsticks semaphores
    int MyLeftChopstickNumber = LeftChopstick(MySeatNumber) ;
    int MyRightChopstickNumber = RightChopstick(MySeatNumber) ;

    // Now get the Win32 'Handles' for those left and right chopstick semaphores and put them into an array
    // this is needed for the Wait_for_multiple_objects call below

    HANDLE MyChopstickHandles[2] = { Chopsticks[MyLeftChopstickNumber]->GetHandle(), Chopsticks[MyRightChopstickNumber]->GetHandle() } ;

    for(int i = 0; i < 10; i++)
    {
        SLEEP(300 + (MySeatNumber * 300)) ;                        // simulate thinking, use a different time delay for each philosopher
        // Now simulate Eating by acquiring chopsticks to left and right of me
        WAIT_FOR_MULTIPLE_OBJECTS(2, MyChopstickHandles, INFINITE) ;
        printf("%d ", MySeatNumber) ;                               // show philosopher is eating by printing his number
        SLEEP(300 + (MySeatNumber * 10)) ;                          // simulate eating, this is a different time delay for each philosopher

        Chopsticks[MyLeftChopstickNumber]->Signal() ;             // put down left chopstick, i.e. release the resource after eating
        Chopsticks[MyRightChopstickNumber]->Signal() ;            // put down right chopstick, i.e. release the resource after eating
    }
    return 0 ;
}
```



Deadlock and Starvation

32

```
int main()
{
    int SeatNumber[NUMBER_OF_PHILOS];           // create an array of seat numbers for the philosophers
    CThread *thePhilosophers[NUMBER_OF_PHILOS]; // create an array of pointers to philosopher threads

    for(int i = 0; i < NUMBER_OF_PHILOS; i++) {
        SeatNumber[i] = i;                       // initialise philosopher seat numbers
        char buff[80];
        sprintf(buff, "Chopstick%d", i);          // generate the name for the mutex based on a number
        Chopsticks[i] = new CMutex(buff);         // create the mutex's using generated name
        thePhilosophers[i] = new CThread(Philosopher, ACTIVE, &SeatNumber[i]); // create the philosopher threads
    }

    for(i = 0; i < NUMBER_OF_PHILOS; i++) {
        thePhilosophers[i]->WaitForThread();      // wait for philosophers to terminate
        delete thePhilosophers[i];               // delete thread object
    }

    for(i = 0; i < NUMBER_OF_PHILOS; i++)         // delete the chopsticks at the end
        delete Chopsticks[i];

    return 0 ;
}
```

- This program can be found in Question 12 of the Tutorial Sheets (Don't worry you don't have to memorize it for an exam)